

# Understanding the Docker Internals



Nitin AGARWAL · [Follow](#)

3 min read · Jan 5, 2017



936



3



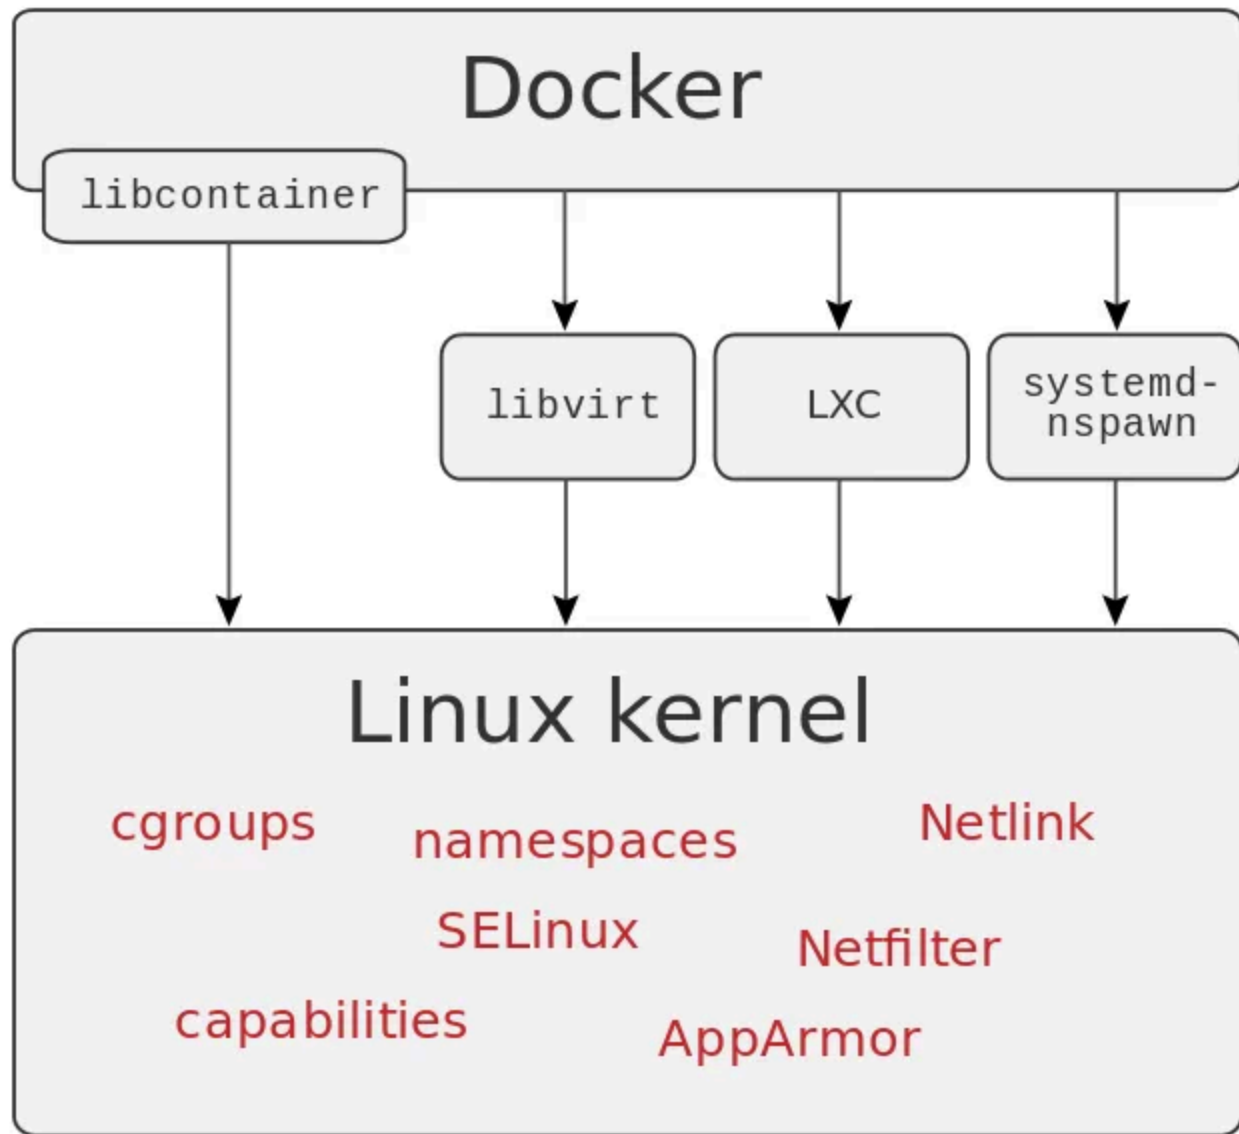


Image Source and Credits: <https://delftswa.github.io/chapters/docker/>

Docker takes advantage of several features of the Linux kernel to deliver its functionality.

## Namespaces

Docker makes use of kernel namespaces to provide the isolated workspace called the *container*. When you run a container, Docker creates a set of *namespaces* for that container. These namespaces provide a layer of isolation. Each aspect of a container runs in a separate namespace and its access is limited to that namespace.

## Docker Engine uses the following namespaces on Linux:

- **PID namespace** for process isolation.
- **NET namespace** for managing network interfaces.
- **IPC namespace** for managing access to IPC resources.
- **MNT namespace** for managing filesystem mount points.
- **UTS namespace** for isolating kernel and version identifiers.

## Docker Grounds up: Namespaces

- Process trees.
- Mounts.
- Network.
- User accounts.
- Hostnames.
- Inter-process communication.

```
pid_t pid = clone(..., flags, ...)
```

CLONE_NEWUTS	hostname, domainname
CLONE_NEWIPC	IPC objects
CLONE_NEWPID	Process IDs
CLONE_NEWNET	Network configuration
CLONE_NEWNS	File system mounts
CLONE_NEWUSER	User and Group IDs

```
setns(int fd, int nstype)
```

```
CLONE_NEWIPC
CLONE_NEWNET
CLONE_NEWUTS
```

```
Also: unshare(flags)
```

Image Source and Credits (Rohit Jnagal): <http://www.slideshare.net/RohitJnagal/docker-internals>

# Docker Grounds up: Processes & Networking

We have resources, isolation, and file system management.

Docker daemon handles starting/stopping processes with:

- Attach logic
- Logs
- TTY management
- Docker run options
- Events and container state

Network Management

- NAT, Bridge, Veth
- Expose
- Links

Image Source and Credits (Rohit Jnagal): <https://docs.google.com/presentation/d/1juVgXiLTM-ZmAmYBOshNwhBABkUqwlxVodHZwq-OeGg/edit?usp=sharing>

## Cgroups

Docker also makes use of kernel control groups for resource allocation and isolation. A cgroup limits an application to a specific set of resources.

Control groups allow Docker Engine to share available hardware resources to containers and optionally enforce limits and constraints.

Docker Engine uses the following cgroups:

- **Memory cgroup** for managing accounting, limits and notifications.
- **HugeTBL cgroup** for accounting usage of huge pages by process group.
- **CPU group** for managing user / system CPU time and usage.
- **CPUSet cgroup** for binding a group to specific CPU. Useful for real time applications and NUMA systems with localized memory per CPU.
- **BlkIO cgroup** for measuring & limiting amount of blkIO by group.

- **net\_cls** and **net\_prio cgroup** for tagging the traffic control.
- **Devices cgroup** for reading / writing access devices.
- **Freezer cgroup** for freezing a group. Useful for cluster batch scheduling, process migration and debugging without affecting prtrace.

## Docker Grounds up: Resource Isolation

### Cgroups : Isolation and accounting

- cpu
- memory
- block i/o
- devices
- network
- numa
- freezer




Image Source and Credits (Rohit Jnagal & Mairin): <http://www.slideshare.net/RohitJnagal/docker-internals> and <https://mairin.wordpress.com/2011/05/13/ideas-for-a-cgroups-ui/>

## Union File Systems

Union file systems operate by creating layers, making them very lightweight and fast. Docker Engine uses UnionFS to provide the building blocks for containers. Docker Engine can use multiple UnionFS variants, including AUFS, btrfs, vfs, and devicemapper.

## Docker Grounds up: Filesystem

### File-system Isolation:

Building a rootfs dir and chroot into it.

With mount namespace, use pivot-root.

### Features:

Layering, CoW, Caching, Diffing

### Solutions:

UnionFS, Snapshotting FS, CoW block devices




Image Source and Credits (Rohit Jnagal): <http://www.slideshare.net/RohitJnagal/docker-internals>

## Docker Grounds up: Filesystem

Open in app ↗

Sign up

Sign in

Medium



Search



Write



	Supercheap	Cheap	Cheap
Changing small files	Superfast Supercheap	Fast Cheap	Fast Costly
Changing large files	Slow (first time) Inefficient (copy-up!)	Fast Cheap	Fast Cheap
Diffing	Superfast	Superfast	Slow
Memory usage	Efficient	Efficient	Inefficient (at high densities)
Drawbacks	Random quirks AUFS not mainline IAUFS more quirks	ZFS not mainline BTRFS not as nice	Higher disk usage Great performance (except diffing)
Bottom line	Ideal for PAAS and high density things	This is the Future (probably)	Dodge Ram 3500

Image Source and Credits (Jerome Petazzoni): <http://www.slideshare.net/RohitJnagal/docker-internals>

### Container Format

Docker Engine combines the namespaces, control groups and UnionFS into a wrapper called a container format. The default container format is

libcontainer.

## Security

Docker Engine makes use of AppArmor, Seccomp, Capabilities kernel features for security purposes.

- **AppArmor** allows to restrict programs capabilities with per-program profiles.
- **Seccomp** used for filtering syscalls issued by a program.
- **Capabilities** for performing permission checks.

### Docker Grounds up: Add Security

- **Linux Capabilities**
  - Drops most capabilities.
  - Enable what a task needs.
- **GRSEC and PAX**
- **SELinux**
- **AppArmor**



Image Source and Credits (Leo Reynolds): <http://www.slideshare.net/RohitJnagal/docker-internals>

. . .

Source: ~ <http://docker-saigon.github.io/post/Docker-Internals/>

Disclaimer: Content and Image source has been mentioned. Special credit to concerned folks.

Docker

DevOps

Cloud Computing

Linux

Cloud



## Written by Nitin AGARWAL


838 Followers

Follow




Senior Software Engineer — Cloud Native and Distributed Systems

### More from Nitin AGARWAL



Nitin AGARWAL



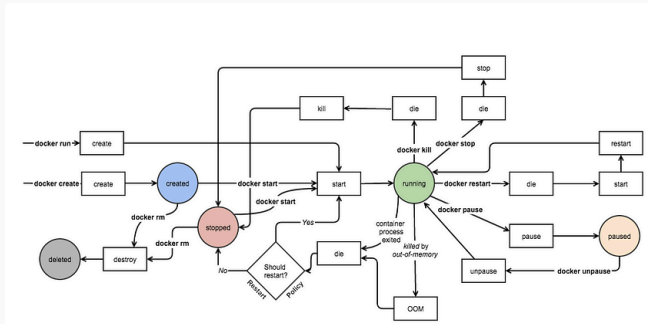
Nitin AGARWAL



## Best Practices for working with Dockerfiles

Dockerfiles allow users to define the exact actions needed to create a new container...

W+



```

graph TD
    1["pid = 1"] --> login["login  
pid = 8415"]
    1 --> kthreadd["kthreadd  
pid = 2"]
    1 --> sshd1["sshd  
pid = 3028"]
    login --> bash["bash  
pid = 8416"]
    bash --> ps["ps  
pid = 9298"]
    bash --> emacs["emacs  
pid = 9214"]
    kthreadd --> khelper["khelper  
pid = 6"]
    kthreadd --> pdflush["pdflush  
pid = 200"]
    sshd1 --> sshd2["sshd  
pid = 3610"]
    sshd2 --> tcsch["tcsch  
pid = 4005"]
  
```

## An Init System inside the Docker Container

In this blog, we'll understand the PID 1 zombie reaping problem and how it can be solved...

W+



## Recommended from Medium

**ALEXANDER NGUYEN**  
Software Development Engineer

Mar. 2020 – May 2021

- Developed Amazon checkout and payment services to handle traffic of 10 Million daily global transactions
- Integrated Iframes for credit cards and bank accounts to secure 80% of all consumer traffic and prevent CSRF, cross-site scripting, and cookie-jacking
- Led Your Transactions implementation for JavaScript front-end framework to showcase consumer transactions and reduce call center costs by \$25 Million
- Recovered Saudi Arabia checkout failure impacting 4000+ customers due to incorrect GET form redirection

**Projects**

**NinjaPrep.io (React)**

- Platform to offer coding problem practice with built in code editor and written + video solutions in React
- Utilized Nginx to reverse proxy IP address on Digital Ocean hosts
- Developed using Styled-Components for 95% CSS styling to ensure proper CSS scoping
- Implemented Docker with Seccomp to safely run user submitted code with < 2.2s runtime

**HeatMap (JavaScript)**

- Visualized Google Takeout location data of location history using Google Maps API and Google Maps heatmap code with React
- Included local file system storage to reliably handle 5mb of location history data
- Implemented Express to include routing between pages and jQuery to parse Google Map and implement heatmap overlay



Alexander Nguyen in Level Up Coding

## The resume that got a software engineer a \$300,000 job at Google.

1-page. Well-formatted.



Jun 1



23K



458



Furkan Türkal

## How does Docker actually work? The Hard Way: A Technical Deep...

Unveiling the power of Docker: What is Docker? How does Docker work? Explore th...

Jun 4



559



6




### Lists




#### General Coding Knowledge

20 stories · 1625 saves



#### Productivity

242 stories · 569 saves



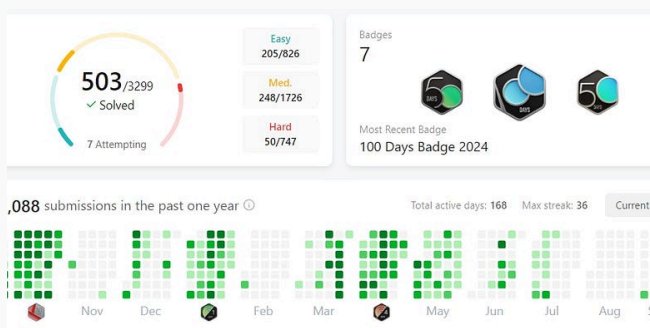
#### Coding & Development

11 stories · 836 saves



#### Staff Picks

745 stories · 1352 saves



Surabhi Gupta in Code Like A Girl




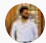
Akhilesh Mishra

## Why 500 LeetCode Problems Changed My Life

How I Prepared for DSA and Secured a Role at Microsoft

★ Sep 26 🖱 1.5K 💬 34



 Sachin Bhatta

## CI/CD Pipeline with GitHub, Jenkins and Docker

A CI/CD pipeline automates code deployment using GitHub, Jenkins, and Docker.

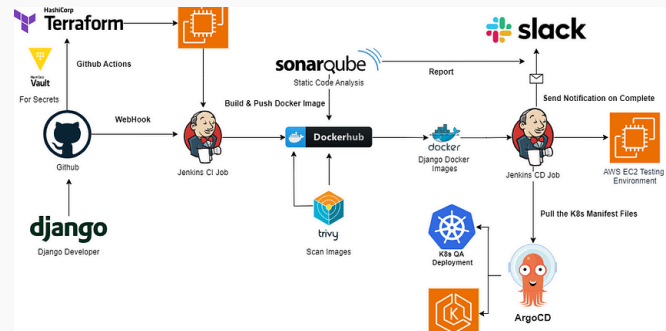
★ Jun 16 🖱 14 💬 1

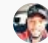


## Everything You Need To Know About Docker

Devops zero to hero #3—A Complete Guide to start using Docker in your devops workflow

★ Jan 23 🖱 522 💬 3



 Joel Wembo in Django Unleashed

## Technical Guide: End-to-End CI/CD DevOps with Jenkins, Docker,...

Building an end-to-end CI/CD pipeline for Django applications using Jenkins, Docker,...

★ Apr 12 🖱 815 💬 16



See more recommendations